

Genetic Mapping of Diseases through Big Data Techniques

Julio C. S. dos Anjos¹, Bruno Reckziegel Filho¹, Junior F. Barros¹
Raffael B. Schemmer¹, Claudio F. R. Geyer¹ and Ursula Matte²

¹*Universidade Federal do Rio Grande do Sul, GPPD, Institute of Informatics,
Caixa Postal 15.064, 91.501-970, Porto Alegre, RS, Brazil*

²*Laboratory of Gene Therapy Hospital de Clínicas de Porto Alegre,
Rua Ramiro Barcelos, 2350, 90035-903, Porto Alegre, RS, Brazil
{jcsanjos, brfilho, jfbarros, raffael.schemmer, geyer}@inf.ufrgs.br, umatte@hcpa.ufrgs.br*

Keywords: Data Intensive Computing, MapReduce, Parallel Distributed Processing, Genome Annotation.

Abstract: The development of sophisticated sequencing machines and DNA techniques has enabled advances to be made in the medical field of genetics research. However, due to the large amount of data that sequencers produce, new methods and programs are required to allow an efficient and rapid analysis of the data. MapReduce is a data-intensive computing model that handles large volumes that are easy to program by means of two basic functions (Map and Reduce). This work introduces GMS, a genetic mapping system that can assist doctors in the clinical diagnosis of patients by conducting an analysis of the genetic mutations contained in their DNA. As a result, the model can offer a good method for analyzing the data generated by sequencers, by providing a scalable system that can handle a large amount of data. The use of several medical databases at the same time makes it possible to determine susceptibilities to diseases through big data analysis mechanisms. The results show scalability and offer a possible diagnosis that can improve the genetic diagnosis with a powerful tool for health professionals.

1 INTRODUCTION

The information provided by DNA sequencing is crucial to ensure the success of biological research. Biotechnology is one of the fields that has benefited the most from this kind of information, especially for the development of new pharmaceutical substances, foodstuffs, pesticides and agricultural products, and when conducting a clinical diagnosis (William J and Palladino, 2012).

DNA sequencing was once a very expensive process, but currently there are techniques and machines that are considerably more cost-effective. The National Human Genome Research Institute (NHGRI) has shown that prices fell between the years 2007 and 2008, when the sequencing methodology migrated from Sanger to what was then called Next Generation Sequencing (NGS). The time taken for sequencing has also increased. As shown in Table 1, a few NGS machines (particularly those that use ion semiconductors such as Ion ProtonTM (Scientific, 2014)) are able to sequence millions of base pairs in a few hours, and hence can produce large amounts of data very quickly. However, at present there is a limited number of tools that can analyze the produced data in an

Table 1: Technical specifications of Ion Torrent sequencing machines.

Characteristics	Machine Models			
	PGM318	PI	PII	PIII
Sensor Number	~11 M	~165 M	~660 M	~1.2 B
Input Size	~2 GB	~10 GB	~32 GB	~64 GB
Execution Time	4~7 hrs	2~4 hrs	2~4 hrs	2~4 hrs
Average Read	400 BP	200 BP	100 BP	100 BP
Number of Reads	~5.5 M	~82 M	~330 M	~660 M
Key: M = Million B = Billion BP = Base Pairs GB = GigaBytes				

automated and rapid fashion.

In this context, MapReduce computing is an attractive solution as it implicitly offers a parallel distributed solution for processing large amounts of data, such as those produced by sequencers. The MapReduce framework is easier to program and has a distributed file system that other parallel systems do not have, *e.g.* MPI. MapReduce improves accuracy and reduces the complexity of biological applications, such as, multiple sequence alignment algorithm (Zou et al., 2014). This means that MapReduce is a framework model that is suitable for this study. In addition, the existing MapReduce implementations, such as Hadoop (an open-source implementation maintained

by Apache Software Foundation) handles as fault-tolerance and performance mechanisms such as load-balancing, automatically.

There is an huge variation in human genomes that can be determined by massive parallel sequencing. However, many of these variations are not clinically relevant. Thus, there is a great need for methods that can discriminate between disease-causing mutations and normal genetic variability in a short run-time (Frebourg, 2014).

The purpose of this work is to develop a genetic mapping system that can assist doctors in the clinical diagnosis of patients by conducting an analysis of the genetic mutations contained in their DNA. The scientific goal is to provide an efficient methodology for the genetic mapping of diseases with big data systems. The execution time of this GMS must be scalable with regard to the amount of processed data. The data is produced by sequencers that employ the ion semiconductor technology. We used the Hadoop MapReduce system and databases from the Gene Report (NCBI, 2014) and Ensembl (Kinsella et al., 2011). This work is available in <<https://github.com/GeneticMapping/GSM.git>>.

The remainder of this work is structured as follows. Section 2 outlines the background that is required for the definition of the problem, the effects of mutations on gene sequences and data required for genetic diagnosis; it also, provides an overview of the MapReduce programming model and the Hadoop platform. Section 3 examines the related work in the literature. Section 4 gives a detailed description of our proposed solution. In Section 5 there is an account of the methodology employed together with an analysis of the obtained results. Section 6 summarizes the conclusions and make recommendations for future work in the field.

2 BACKGROUND

This section provides background information on genetics and sequencing and their application in medicine. It also examines the technologies used in these areas and the problems currently faced in finding a suitable means of analyzing the impact of mutations on the genome sequencing of patients.

2.1 Nucleotides, DNA, RNA and Genome

Nucleotides are biological molecules that compose the structure of nucleic acids (Alberts et al., 2014). They are formed from a purine base (adenine or

guanine) or a pyrimidine base (thymine, cytosine or uracil). Nucleotides are the basis for forming nucleic acids that are large biological molecules, and essential to all known forms of life. Deoxyribonucleic Acid (DNA) contains the genetic code of a living being, which comprises long sequences of nucleotides that form a double helix structure. The Ribonucleic Acid (RNA), which is often formed by a simple chain, synthesizes the proteins from a cell on the basis of genetic information contained in the DNA and conveyed in the Messenger RNA (mRNA). The basic differences between the DNA and RNA structures lies in the number of helices each has (the former has two whereas the latter has only one), and in the change of a base from thymine to uracil (William J and Palladino, 2012).

The DNA physical structure has a three dimensional shape, which is the result of the forces exerted by the electrons that compose the bases. Hence, any change in one of the bases, alters the position of the helix and its shape, and turns its function into a biological system. The change in molecular structure can lead to a genetic mutation that may, or may not, be associated with the propensity an individual has to contract a certain disease (Sawyer et al., 2007).

By complying with the rules laid down by the genetic code, the genetic information provided by the DNA sequences via mRNA is translated into an amino acid sequence, thus generating a protein (William J and Palladino, 2012). In this process, as shown in Figure 1 which is adapted from (William J and Palladino, 2012), each triplet of nucleotides is called a *codon*, and each *codon* represents an amino acid, that synthesizes a protein. However, only the genomic regions defined as coding regions are considered in the process of protein translation.




Figure 1: Central dogma of molecular biology.

The genome represents all the genetic information inherited by an organism. The human genome consists of a large amount of DNA divided into organized

structures called chromosomes. The human genome contains 23 chromosome pairs and an estimated number of 32,000 genes. The total number of DNA bases is around 3.6 billion.

2.2 Mutations, Polymorphism and Clinical Genetics

A mutation is defined as a change in the nucleotide sequence of an organism. Mutations can be caused by the irreparable damage suffered by the genome, errors in the replication process or the insertion/deletion of DNA fragments by mobile genetic elements. Several studies, such as those carried out by (Johnsen et al., 2013), suggest that, if a mutation changes the protein that a gene produces, the end result will likely be harmful to the organism. The coding region of a genetic code begins with a start *codon* and ends with a stop *codon*; this region is the focus of our analysis. However, it should also be noted that, some mutations do not modify the amino acid generated (i.e., even if a mutation takes place, the generated protein remains the same).

Polymorphism is a kind of mutation that takes place with a frequency greater than 1% in a population, and can be divided into distinct, well-defined classes. An example is the groups of the ABO blood group system (classes A, B, AB and O). According to (Nussbaum et al., 2013), there are three classes of mutations, namely those that affect the number of chromosomes in a cell (also called genomic mutations), those that alter the structure of specific chromosomes (chromosomal mutations), and mutations that change individual genes (gene mutation). This work focuses on conducting an analysis of gene mutation; however, the impact of other mutations on the health of patients will be addressed in future work.

2.3 MapReduce

MapReduce is a programming framework that abstracts the complexity of parallel applications by partitioning and scattering data sets across hundreds or thousands of machines, and by bringing computation and data closer (Dean and Ghemawat, 2010). The Figure 2, adapted from (White, 2012), shows the MR data flow. The Map and Reduce phases are handled by the programmer, whereas the Shuffle phase is created while the job is being carried out. The input data is split into smaller pieces called chunks, that normally have a size of 64 MB. The data is serialized and distributed across machines that compose the Distributed File System (DFS).

When running an application, the job is split by the master into several Map and Reduce tasks; following this, it assigns tasks to workers that then run each processing stage. The machine that is given a Map task, handles a Map function and emits key/value pairs as intermediate results that are temporarily stored in the worker disks. The execution model creates a computational barrier, which allows tasks to be synchronized between the Map and Reduce. A Reduce task does not start its processing until all the Map tasks have been completed.

A hash function is applied to the intermediate data produced during the Map phase to determine which keys will compose a Reduce task. All the pairs combined with these keys are transferred to one machine, during the Shuffle, so that they can be processed by a Reduce task. After a reduction function has been applied to this data, a new key/value pair is issued. The result is then stored in the distributed file system and thus can be made available to the client who submitted the job. The Shuffle phase consists of two stages: one performed on the machine that processes the Map task, which sorts the keys and serializes the data. The other is performed after the intermediate data has been sent to the reducer machine, which arranges the received data to allow the keys to be properly grouped and then runs the Reduce task. The Hadoop implements the Combiner function to save bandwidth. The Combiner is a preprocessing of intermediate keys into memory of workers during the map phase to form the output for input data in the Reduce phase. (White, 2012).

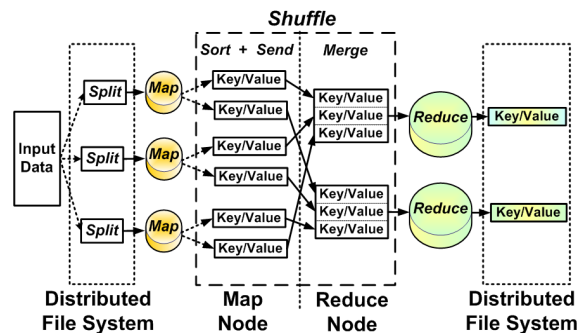


Figure 2: MapReduce data flowchart.

The management mechanisms, data replication and execution control were added to the framework during Hadoop implementation. The management architecture is based on the master/worker model, while a slave-to-slave data exchange requires a P2P model (White, 2012). Hadoop MapReduce is an implementation for a wide range of Big Data applications used by large companies like Facebook, Amazon, Cloudera and IBM. Most works, (including this system) are

based on a Hadoop implementation which is regarded as the most advanced open-source implementation of MapReduce (Dean and Ghemawat, 2010).

3 RELATED WORKS

The advanced technique provides a collective solution for the problems currently facing mankind and the NGS genome. It is a challenge to data processing, in terms of computational resources (O'Driscoll et al., 2014). When dealing with big data, it is necessary to build new models or adjust existing ones by taking account of the data dispersion and processing capability. In establishing the relationship between MapReduce and genomic analysis, it is necessary to understand the characteristics that make the model important for this research field. When supported by the MapReduce model, the application development is particularly well suited to common “*map-reduce scan*” pipelines. Schatz introduced some important tasks for application development that can be achieved with MapReduce (Schatz, 2010), such as the following:

1. Map: Genomic problems which are composed of many sequences that can be mapped to the reference genome in parallel with multiple machines.
2. Reduce: Sequences of alignments are aggregated so that all the alignments on the same chromosome or locus are grouped together and sorted by position.
3. Scan: The sorted alignments are scanned to identify biological events, such as polymorphism or different expressions.

On the basis of the above, it is worth discussing the open MapReduce framework (Apache Hadoop) that is used to build apps that conform to the needs of genomics. In addition, there are libraries such as Hadoop-BAM (Niemenmaa et al., 2012) that are used to dealing with BAM files (a compressed file format that is used for NGS and usually well suited to processing genomic sequences in a large number of machines). Crossbow (Gurtowski et al., 2012) is another example of how Hadoop can be used to discover Single Nucleotide Polymorphism (SNP) from sequences of data in cloud systems or Hadoop clusters. In a similar way, our work is able to find gene mutations at the first moment and afterwards, when making comparisons with databases, may detect possible diseases.

In the genomic mapping of sequences, CloudBurst (Schatz et al., 2010) and ClodAligner (Nguyen et al., 2011) are both cases where Hadoop can be used for

data analysis based on NGS sequencers for discovering SNP sequences, as well as to build alignment for short-reads or long-reads of genomic data in reference to a human genome. It is worth taking note of the capability for scalability in these works through cluster and cloud, as is shown in our work.

From the perspective of adopting an interface that is user-friendly and easy for users, there are a range of approaches that use MapReduce, such as the Genome Analysis Toolkit (GATK) (McKenna et al., 2010) where the aim is to create a functional language control flow by supporting the development of an NGS application program. When compared with a traditional programming structure, it requires less time for development. Another example is CloudDOE (Chung et al., 2014) that frees scientists from the need to carry out complicated procedures. Our study makes possible to create a system that saves time for user when are genomic problems are being addressed.

In the case of health-care, Hansen stresses the benefits that can be derived, (such as Big Data), from understanding factors such as the prevention of diseases, detection of modifiable risk factors for disease and intervention to bring about behavior change (Hansen et al., 2014). As mentioned earlier, accessibility is an important point to take into account. Our study checks several disease databases to find a possible relation of diseases with genomic mutations, with a clear and simple method, based on the data output of the Clinical Hospital of Porto Alegre (HCPA) scientists. This study has the potential to help to find mutations that are able to generate pathologies prior to their action on an individual.

3.1 Genome Annotation

According to Costa, it is difficult to find an equivalent system to genome annotation (Costa, 2014). There are few Big Data projects that are aimed at supporting the clinical area. The author cites Appistry Cloud, as a cloud computing system for genetic analysis which is concerned with oncology research and the development of new drugs. Unlike our implementation, this seeks to provide a tool to annotate generic pathologies. CLC Bio is another work cited by Costa, as a system for diagnosing cancer pathologies, although this requires the user's computer to install an application to insert data and create scripts. These approaches are very complex for users without experience, in contrast with GMS, where the user only needs to execute a web application.

The DNAnexus project (BCM, 2014) provides solutions for DNA sequencing centers using NGS sequencers, since the application is developed for a

Cloud platform, but with an analysis and sequencing linked to DNA alignment. Although the analysis does not involve polymorphism detection. The NGS project (McKenna et al., 2010) shows a toolkit implementation in a MapReduce framework for an exome analysis, the goal of which is capture sequencing, and unlike other studies, our implementation is aimed at identifying significant genetic mutations from a gene.

4 DESCRIPTION OF THE MODEL

Currently there are only a few applications for clinical diagnosis that provide an automated genetic analysis of patients with high scale of amount of data generated by NGS. A manual method was previously used from the clinical diagnosis of patients by researchers at the Clinical Hospital of Porto Alegre (HCPA), in Brazil. Genetic data for each research study was obtained from sequencers and entered by hand into Excel spreadsheets. After the data had been allocated correctly, they had to be separated into *codons* and compared with a mutation list that included the known diseases. This process is time-consuming.

The Ion Torrent sequencers can generate data in multiple formats, such as BAM, VCF and FASTA (MEDLINE, 2013). The databases that provide references for research are widely dispersed among a hierarchy of classes. Several of these databases are maintained by Government bodies such as the National Center for Biotechnology Information (NCBI): the database that stores information about several specific genetic mutations (dbSNP); the database that references to the dbSNP and provides information about mutations (Ensembl); the database that keeps information about the composition of several coding areas of genes (CCDS) and the Gene Report, a database that keeps information about mutations that are considered pathogenic, directly related to diseases. The input data format used for these studies is similar to FASTA.

This work implements a code by using the MR programming model of Hadoop to analyze and process the amount of data generated by the sequencing process. Before starting the MR processing, the input data is previously preprocessed by python scripts. This preprocessing step indexes a cluster of genomic analysis from several patients and forms a single input data file. Each line of this file is composed by one gene from one patient. The line is tagged by the patient identifier and the name of the gene. These data are used in MR processing to determine how the gene reference base should be used in the processing. This model assumes that the input file is formed by differ-

ent kinds of genes and patients.

Figure 3 shows an overview of the flowchart of this proposal. The flowchart is divided into three stages. The first stage is responsible for persisting the input data composed by the patients' gene in the Hadoop distributed file system (HDFS), where the data is split into chunks. Each chunk consists of several lines of the input data, where each line corresponds to one gene from each patient. In the second stage, a preprocessing of the input file is carried out to find gene's list from all patients. One Python script search in CCDS database to find a reference gene code, and index this information in a local memory cache of each worker node. The third stage is to compose the annotation by using the code that implements the MR processing flow. The code consists of two methods.

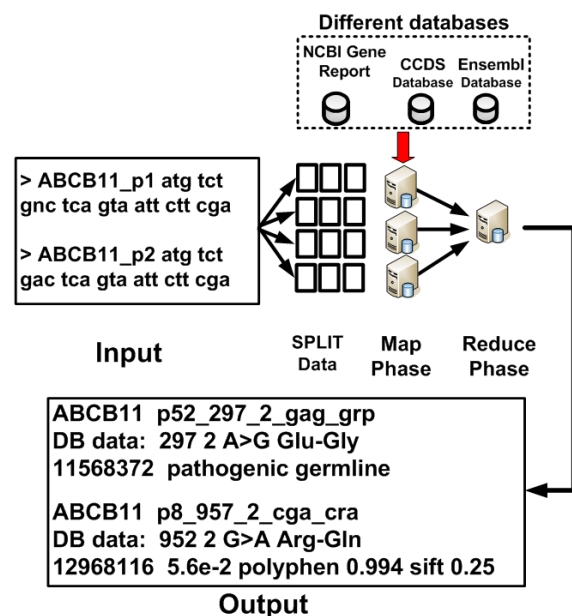


Figure 3: Flowchart of the proposal.

The first method is calling Map and is responsible for reading the input data formed by the gene analysis in patients and to compare the data with the reference gene. The Map function notes the position where the difference is found, the reference tuple of the gene and the input tuple of the patient gene. When a difference is found in the input, this is compared with a reference of genetic mutation in both Ensembl and Gene Report databases. This comparison executes into Combiner on Map phase.

The Combiner receives this data and searches in the mutation databases. The search uses the position and the name of the gene to query the databases and writes the pathology data found. If one pathology is reported, the Map emits a key/value pair intermedi-

ate. The key is formed by the position in the gene and the value within a tuple of a reference gene and the patients' gene.

After, the Reduce function emits a new key/value pair with the information about any associated pathologies, for each patient and saves on HDFS. Only mutations found in the databases are written in the output, followed by any messages from one associated pathology to this mutation.

5 EVALUATION

This section describes the environment setup and results of the evaluation as a means of demonstrating the features and scalability of our proposal.

5.1 The Environment Setup

Two environments have been considered. The first is a cluster machine comprising 19 nodes with a heterogeneous configuration. In more detail, the cluster is composed of 5 machines with a P4 2.79GHz processor, 2GB of RAM, and 400GB of hard disk; 14 machines with an Intel P4 2.79GHz Hyper Thread processor, 2GB of RAM and 1TB hard disk; and a master node with a QuadCore 3.1GHz processor, 4GB of RAM and 1TB hard disk. The system is the Hadoop version 1.0.4. The second is a cluster in Cloud Microsoft Azure that comprises 19 nodes, A1 type. This node comprises a single core of Intel Xeon E5-2660, 1.75GB of memory RAM and 30GB of HDD. All of the nodes have the same configuration. This study relies on the same version of Hadoop and Linux kernel in Azure in an attempt to approximate the results. Azure settings were used to ensure that all the nodes were allocated in the same geographic location.

In the experiments, we vary the number of used worker nodes (*i.e.* 1, 2, 5, 10, 15 and 19) and the input data size (*i.e.* 2GB, 10GB and 20GB). The number of executions generated 18 different cases and each case was executed 30 times, thus resulting in 540 experiments being performed in the cluster. In all the tests, the block size used by Hadoop for the data replication and processing was 64MB and the replication factor was set to 3. The standard deviation was lower than 5%.

5.2 Results and Analysis

The experiment introduced by Figure 4 and Figure 5 shows the GMS execution time of the cluster and the cloud. In the same configuration, it is often observed that as the data input size increases, the execution time

increases proportionally in the same way. The execution time is ≈ 2 hours for one standalone machine (Intel P4 2.79GHz, 4GB RAM, HD 500GB).




Figure 4: GMS execution time in the GradeP cluster.

Similarly, when more nodes were added to allow the processing task to be carried out in a parallel way, the time decreases proportionally to an increase in the number of nodes. This occurs because the data-intensive applications, developed by means of Hadoop and the MR programming model, have a greater scalability and better performance when there is a larger volume of data.




Figure 5: GMS execution time in Cloud (Microsoft Azure).

The results showed that even the cloud configuration has more powerful resources like newer CPU cores than those of the cluster, the execution time is slower than with the cluster. This can be explained in function of shared network overhead on cloud. In addition, Azure has to control the state of the virtual machine that increases this overhead. However, both of the configurations are scalable, which reduces the run-time of all the input sizes in all the variations of the node numbers. This shows that the solution pro-

posed by our work can be used either in cloud or cluster environments, because it is scalable when there is a large configuration of machines, and since the speedup is considerable, it sets out from a sequential time provided by one node. In addition, the increase of the input size in the DNA sequencing increases the requirements and the problem of the processing can be resolved by adopting these approaches and techniques.

Figures 6 and 7 show the speedup. This metric represents the acceleration of sequential execution time when more than one node is used. The ideal speedup is always the number of nodes used in the execution.




Figure 6: GMS speedup in a Cluster.




Figure 7: GMS speedup in Cloud (Microsoft Azure).

The results in both cluster and cloud show that the increase of speedup is related to the size of the input data. In all the scenarios, the speedup improves when more data are employed in the processing; this behavior is linked to MapReduce execution in Big Data. The overhead required in Cloud for managing the resources causes wait for the data to be processed.

The size of the input data can cover this time.

These results prove that large distributed machines (ranging from a dozen to thousands of computed nodes) correspond to the size of the data generated by NGS sequencers, like Ion Torrent for example. The last and most important point that needs to be made is the question of achieving a better speedup by cluster than by cloud. This performance is linked to the execution time and the overhead (related by cloud) which involves managing costs and sharing resources through a network connection with the nodes.

5.2.1 Output Analyzis

The Figure 8 shows one example from the GMS output. The detection of mutation A>G Glu-Gly (pathogenic germline) in Position 2 of gene 297 compared with reference DNA indicates the susceptibility of colorectal cancer. The detection of mutation G>A Arg-Gln in Position 2 of gene 957 compared with reference DNA, indicates the susceptibility to gallbladder cancer. This information demonstrates the correct framework execution.

```

ABCB11 p52_297_2_gag_grp
DB data: 297 2 A>G Glu-Gly
11568372 pathogenic germline

ABCB11 p8_957_2_cga_cra
DB data: 952 2 G>A Arg-Gln
12968116 5.6e-2 polyphen 0.994 sift 0.25

```

Figure 8: GMS output.

6 CONCLUSION

This work proposes and implements GMS to analyze and compute large amounts of data generated by NGS sequencers such as Ion Torrent. In addition, this study is a part of a research project that seeks to automate the process of collecting and processing of patient gene analysis, using techniques like a Big Data framework and distributed environments to overcome the limitations of traditional technologies, related to a single personal computer.

The proposed GMS addressed the problems arising from automation and the handling of the data. The detection of mutations *e.g.* A>G in Position 2 of gene 297 compared with reference DNA and G>A in Position 2 of gene 957 compared with reference DNA, indicating the susceptibility of determined cancer type, demonstrates that our GMS proposal achieves its initial goal. This work makes use of two scenarios involving distributed machines, a cluster and cloud computing. In the same scenarios, the application proposes to enhance performance when

using a greater number of resources. The results suggest that both cluster and cloud can be used to achieve a reduction of run-time, but the cloud scenery have a overhead improved.

In future work, we believe that the researches should to improve the features supported by the proposed, such as those that concern generating feedback from a mutation found in a patient gene to a database of mutations like Ensembl or Gene Report can be implemented. Other future study could attempt to implement this system, for example, to finding an online user-friendly solution in cloud. However, many issues need still to be discussed to be explored graphic interfaces for the use of cloud to persist gene analysis of the patients.

ACKNOWLEDGEMENTS

This work was supported by several institutions, including the following: CNPq through grant number MCT/CNPq N^o 70/2009 – PGAEST- MCT/CNPq; FAPERGS via grant FAPERGS/CNPq 008/2009. It was also assisted by the research project “GREEN-GRID: Sustainable HPC”; Grid5000, a grid platform developed by the INRIA ALADDIN project, and the support provided by CNRS, RENATER and other universities. Some experiments were assisted by Microsoft Azure environment from Microsoft Research. The research was also partly sponsored by CAPES grant 99999.014966/2013-01 (through DSE program) and by CNPQ-PIBITI-UFRGS FUNTTEL - MAREMOTO project. Members from the GPPD MapReduce group also assisted in the development of this project.

REFERENCES

- Alberts, B., Johnson, A., Lewis, J., Raff, M., Roberts, K., and Walter, P. (2014). *Molecular Biology of the Cell*. Garland Science, 6th edition.
- BCM (2014). DNA Nexus Project. Technical report.
- Chung, W.-C., Chen, C.-C., Ho, J.-M., Lin, C.-Y., Hsu, W.-L., Wang, Y.-C., Lee, D. T., Lai, F., Huang, C.-W., and Chang, Y.-J. (2014). CloudDOE: A User-Friendly Tool for Deploying Hadoop Clouds and Analyzing High-Throughput Sequencing Data with MapReduce. *PLOS ONE*, 9:e98146.
- Costa, F. F. (2014). Big data in biomedicine. *Drug Discovery Today*, 19(4):433–440.
- Dean, J. and Ghemawat, S. (2010). MapReduce - A Flexible Data Processing Tool. *Communications of the ACM*, 53(1):72–77.
- Frebourg, T. (2014). The challenge for the next generation of medical geneticists. *Hum Mutat*, 35(8):909–11.
- Gurtowski, J., Schatz, M. C., and Langmead, B. (2012). Genotyping in the cloud with Crossbow. *Curr Protoc Bioinformatics*.
- Hansen, M., Miron-Shatz, T., Lau, A. Y. S., and Paton, C. (2014). Big Data in Science and Healthcare: A Review of Recent Literature and Perspectives. *Yearbook of medical informatics*, 9(4):21–6.
- Johnsen, J. M., Nickerson, D. A., and Reiner, A. P. (2013). Massively parallel sequencing: the new frontier of hematologic genomics. *Blood*, 122(19):3268–3275.
- Kinsella, R. J., Kahari, A., Haider, S., Zamora, J., Proctor, G., Spudich, G., Almeida-King, J., Staines, D., Derwent, P., Kerhornou, A., Kersey, P., and Flicek, P. (2011). Ensembl BioMart: a hub for data retrieval across taxonomic space. *Database*, 2011:1–9.
- McKenna, A., Hanna, M., Banks, E., Sivachenko, A., Cibulskis, K., Kernysky, A., Garimella, K., Altshuler, D., Gabriel, S., Daly, M., and DePristo, M. A. (2010). The Genome Analysis Toolkit: A MapReduce framework for analyzing next-generation DNA sequencing data. *Genome Research*, 20(9):1297–1303.
- MEDLINE (2013). *The NCBI Handbook*, volume NBK143764. National Center for Biotechnology Information, 2nd edition.
- NCBI (2014). A Base Pathogenic Mutations. Technical report.
- Nguyen, T., Shi, W., and Shi, W. (2011). CloudAligner: A fast and full-featured MapReduce based tool for sequence mapping. *BMC Research Notes*, 4(171):1–16.
- Niemenmaa, M., Kallio, A., Schumacher, A., Klemela, P., Korpelainen, E., and Heljanko, K. (2012). Hadoop-BAM: directly manipulating next generation sequencing data in the cloud. *Bioinformatics*, 28(6):876–877.
- Nussbaum, R., McInnes, R., and Willard, H. (2013). *Thompson Genetics in Medicine*. Elsevier Science Publishers B. V., 7th edition.
- O’Driscoll, A., Dugelaite, J., and Sleator, R. D. (2014). Big data, Hadoop and cloud computing in genomics. *Journal of Biomedical Informatics*, 46(5):774–781.
- Sawyer, S. A., Parsch, J., Zhang, Z., and Hartl, D. L. (2007). Prevalence of positive selection among nearly neutral amino acid replacements in Drosophila. *Proceedings of the National Academy of Sciences*, 104(16):6504–6510.
- Schatz, M. C., Langmead, B., and Salzberg, S. L. (2010). Cloud computing and the DNA data race. *NATURE BIOTECHNOLOGY*, 28(7):691–693.
- Scientific, T. F. (2014). Choose Next-Generation Sequencing or Sanger Sequencing Solutions. Technical report.
- White, T. (2012). *Hadoop - The Definitive Guide*, volume 1. O’Reilly Media, Inc., 3rd edition.
- William J, T. and Palladino, M. A. (2012). *Introduction to Biotechnology*, volume 1. Pearson, 3rd edition.
- Zou, Q., Li, X.-B., Jiang, W.-R., Lin, Z.-Y., Li, G.-L., and Chen, K. (2014). Survey of MapReduce frame operation in bioinformatics. *Briefings in Bioinformatics*, 15(4):637–647.